

The Air Canada Chatbot Liability Ruling (2024)

Suresh L. Paul

2026-03-11

Introduction

In February 2024, a Canadian civil tribunal made a ruling that sent ripples through the business and technology world. A man named Jake Moffatt had relied on information given to him by Air Canada's customer service chatbot to book a bereavement fare — a discounted airline ticket offered to people traveling due to a family member's death. The chatbot gave him wrong information. Air Canada refused to honor it. And a Canadian court ruled that **Air Canada was legally responsible** for what its own chatbot told a customer.

The case is small in financial terms — the disputed amount was a few hundred Canadian dollars. But its implications are enormous. It raises fundamental questions about **who is accountable when an AI system gives someone bad advice**, and what it means for organizations to deploy AI in customer-facing roles.

Background: Bereavement Fares and the Chatbot

What Is a Bereavement Fare?

Airlines sometimes offer **bereavement fares** — specially reduced ticket prices for people who need to travel urgently because a close family member has died or is dying. These fares acknowledge that grieving people often cannot plan ahead, and that last-minute flights are typically very expensive.

Different airlines have different policies about bereavement fares: who qualifies, how much the discount is, how to apply, and crucially — **whether you can request a reimbursement after the fact** (i.e., buy a regular ticket during an emergency and then later apply for a bereavement fare retroactively).

What Happened to Jake Moffatt

In November 2022, Jake Moffatt's grandmother passed away. He needed to book last-minute flights between Vancouver and Toronto. Before booking, he consulted Air Canada's website chatbot to understand his options.

The chatbot told him that he could book his tickets at the regular price and then **apply for a bereavement fare discount within 90 days of travel** — essentially, buy now and request a refund later. Moffatt took screenshots of this conversation, booked his tickets at full price, and later filed a claim with Air Canada for the bereavement fare discount.

Air Canada rejected the claim. The airline's actual bereavement fare policy required customers to request the fare **at the time of booking**, not after the fact. The chatbot had given him **incorrect information**.

Air Canada’s response was remarkable in its framing: the airline essentially argued that the chatbot was a **“separate legal entity”** from Air Canada and that the company was not responsible for what the chatbot said. The airline also pointed out that its correct bereavement fare policy was available elsewhere on its website.

The Tribunal’s Ruling

The **British Columbia Civil Resolution Tribunal** rejected Air Canada’s arguments entirely.

The tribunal ruled that:

1. **Air Canada is responsible for all information on its website**, including information provided by its chatbot. The company cannot disclaim responsibility for a tool it built and deployed.
2. The argument that the chatbot was a “separate entity” was without merit. A company cannot deploy software to interact with customers and then deny accountability for what that software says.
3. The fact that correct information existed *elsewhere* on the website did not absolve Air Canada. The company should have ensured its chatbot provided accurate information.
4. Air Canada had **a duty to ensure its chatbot was accurate**, and it failed to meet that duty.

The tribunal ordered Air Canada to reimburse Moffatt the difference between what he paid and what the bereavement fare would have been — approximately **CAD \$650** — plus interest and fees.

Why This Case Matters Far Beyond \$650

The dollar amount is trivial. The principle is not.

Organizations Cannot Hide Behind Their AI

Before this ruling, some companies deployed chatbots with the implicit hope that if the chatbot gave incorrect information, the company could blame the software and escape liability. This ruling makes clear that **courts will not accept that argument**. If you deploy an AI system to represent your company to customers, you own what it says.

This is consistent with how we think about human representatives. If a call center employee gives a customer wrong information and the customer relies on it to their detriment, the company is generally liable — not the employee personally. The same logic now applies to AI.

The “Information Available Elsewhere” Defense Does Not Work

Air Canada argued that since the correct policy was posted elsewhere on its website, Moffatt should have found it and disregarded the chatbot. The tribunal rejected this.

This matters enormously for companies deploying AI. Customers trust the sources they interact with. If a company puts a chatbot on its website and a customer asks the chatbot a question, the customer reasonably expects the answer to be accurate. Pointing to a buried policy page as a defense is not acceptable.

Accuracy Is a Non-Negotiable Requirement

Chatbots and AI assistants are built using large language models — sophisticated AI systems that can generate fluent, confident-sounding text on virtually any topic. But fluency and confidence are not the same as accuracy. These systems can and do **make things up** — a problem sometimes called **hallucination** in AI terminology.

Hallucination in AI refers to when an AI system generates information that sounds plausible but is factually incorrect. This is not the AI “lying” — it has no concept of truth or falsehood. It is generating text based on patterns, and sometimes those patterns lead to incorrect outputs.

Deploying an AI that can hallucinate in a customer-facing role — where customers will rely on the information to make real decisions — is a significant liability risk, as this case demonstrates.

Broader Implications: AI Governance

This case is part of a larger emerging conversation about **AI governance** — the policies, oversight mechanisms, and accountability structures that determine how AI systems are built, deployed, and managed.

What Is AI Governance?

AI governance asks questions like:

- Who is responsible when an AI system causes harm?
- How do we ensure AI systems are accurate and fair?
- What disclosures should users receive about AI’s limitations?
- What mechanisms exist for recourse when AI gives bad advice?

Before cases like *Moffatt v. Air Canada*, these questions were largely theoretical. Now they are being answered by courts, regulators, and legislators around the world — and organizations need to be prepared.

The European Union’s AI Act

The European Union passed the **AI Act** in 2024 — the world’s first comprehensive legal framework governing artificial intelligence. It classifies AI applications by risk level and imposes requirements accordingly. Customer-facing AI in sectors like financial services, healthcare, and critical infrastructure faces the strictest requirements, including transparency obligations and human oversight requirements.

Similar regulatory activity is underway in the United States, Canada, the United Kingdom, and other jurisdictions.

What Organizations Should Be Doing

The *Air Canada* case suggests several practical steps for any organization deploying AI in customer-facing roles:

1. **Test AI outputs rigorously** — particularly for high-stakes information like pricing, policies, legal rights, and health guidance. If the AI gets those wrong, the organization bears the consequences.

2. **Build in human oversight** — ensure there is a mechanism for customers to escalate to a human and for humans to review and correct AI-generated information.
3. **Disclose AI limitations clearly** — let customers know they are interacting with an AI, and provide clear guidance about verifying important information.
4. **Keep AI knowledge current** — company policies change. If a chatbot is trained on old information and gives outdated answers, the company may be liable for the difference.
5. **Have a resolution process** — when AI gives wrong information and a customer relies on it, have a clear, fair process for making it right. Air Canada’s refusal to honor the chatbot’s advice was both legally unsuccessful and a public relations disaster.

The Ethical Dimension

Beyond the legal question, this case raises an ethical one. Jake Moffatt was grieving the loss of his grandmother. He was trying to navigate a stressful, emotionally difficult situation and sought help from Air Canada’s systems. He did everything right — he asked a question before booking, he documented the answer, and he followed up through official channels.

Air Canada’s initial response was to blame the software and deny responsibility. That response — even if it had been legally defensible — reflects poorly on the company’s values. Organizations deploying AI in customer-facing contexts should ask not just “are we legally protected?” but “are we treating our customers fairly?”

Key Concepts for Discussion

Concept	Definition	Relevance to This Case
AI Hallucination	When AI generates confident but incorrect information	The chatbot gave factually wrong information about Air Canada’s policy
AI Governance	Policies and accountability structures for AI systems	The ruling establishes that organizations are accountable for their AI
Vicarious Liability	Being held responsible for the actions of an agent acting on your behalf	Air Canada was held liable for its chatbot’s statements
Large Language Model (LLM)	The AI technology underlying most modern chatbots	Powers chatbots but cannot guarantee accuracy
Customer Experience (CX)	The overall experience of interacting with a company	The case highlights the risk of poor AI design to customer trust

Analysis Questions

1. Describe the company featured in the case, including its competitive advantage (or “moat”), historical background, and the nature of its industry or business operations.
2. Summarize the key facts presented in the case. Do these facts relate to or mirror any other cases discussed in class? If so, what insights or lessons from those previous cases can be applied here?
3. Analyze the economic consequences of the information system failure described. If a similar failure occurred in today’s technological environment, would the impact be greater or smaller? How likely is such an incident to occur now?
4. Discuss the broader implications of the risks identified in the case for contemporary IT infrastructures, cloud computing services, and emerging artificial intelligence systems.
5. Identify the key human decision point in this case — the moment where a different choice by a person or team would most likely have prevented or significantly reduced the harm. Who was responsible, and what should they have done differently?
6. Identify and explain five key lessons or takeaways derived from this case study.

Discussion Questions

1. What does this case reveal about where organizations tend to underestimate risk — in their technology, their processes, or their people? Use specific evidence from the case to support your answer.
2. If you were the CIO of the organization in this case, what are the first three changes you would make after this incident, and why?
3. Given that large language models cannot guarantee factual accuracy, under what conditions — if any — is it responsible to deploy them in customer-facing roles? Are there categories of questions where the risk is acceptable and others where it is not?
4. Air Canada argued its chatbot was a “separate entity” and therefore the company was not responsible for its output. Why did the tribunal reject this argument, and do you agree with the ruling?
5. What processes should any organization have in place before deploying a customer-facing AI system — particularly for high-stakes information like pricing, policies, or legal rights?
6. Beyond legal liability, what are the ethical obligations of a company that deploys AI to interact with customers in emotionally difficult situations? Does the context of Moffatt’s bereavement change how you think about Air Canada’s response?

Suggested Further Reading

- Moffatt v. Air Canada, [2024 BCCRT 149](#)
- “Air Canada Must Pay for Chatbot’s Bad Advice on Bereavement Fares” — [The Guardian](#), February 2024
- “The Problem of AI Hallucination” — [MIT Technology Review](#)

Instructions

1. Use **Bloomberg News** as the primary source for case research. You may also reference other *reputable* public websites or databases. All sources must be properly cited in **APA format** within your written report.
2. The **written case report** should be between **5 and 10 double-sided pages**, inclusive of tables, figures, infographics, or any analytical materials. Acceptable submission formats include **Canva, Microsoft Word, PDF, Markdown, or LaTeX**.
3. The **case presentation** should consist of **10 to 30 slides**, incorporating tables, charts, infographics, or analytical visuals as needed. Acceptable file formats include **Canva, Microsoft PowerPoint, PDF, Quarto Markdown, or LaTeX**.
4. Include a **detailed log of each team member's contributions** as an appendix to the written report. Every team member must be prepared to address questions during the presentation, both on the overall case content and specifically on the sections they contributed to.
5. **AI tools may not be used** to generate written content for the report. Collaborative work through **Google Docs** (with version control and simultaneous editing) is encouraged to demonstrate authentic team contribution and transparency in authorship. All version control tracking links must be shared with the instructor when asked.

*Prepared for Suresh Paul's MIS Course — Case Study Student Reading. Not intended for commercial use.
Not intended for AI training*